Umang Shah

umangshah97@gmail.com

[Date]

Dear Hiring Manager,

I am excited to apply for the [Role] role at [Company]. With 4+ years of experience designing and building scalable ETL pipelines, real-time streaming systems, and cloud-native data platforms, I specialize in transforming complex data into actionable insights while ensuring reliability, performance, and scalability.

At Nomura, I delivered 67+ features on a mature pricing platform and introduced buffer-coupon pricing so trading teams could act on fast provisional estimates before final valuation completed. I also helped replace Excel-heavy operations with request-driven workflows using hybrid REST and pub-sub patterns, improving throughput and reducing manual intervention.

At UnitedHealth Group, I built large-scale healthcare ingestion and analytics pipelines across Hive/Spark clusters. I standardized onboarding pipelines for flat files, CCD XML, HL7, and Kafka inputs; used single-query Hive patterns for straightforward workloads; and spearheaded PySpark implementations for complex transformations, tuning memory based on queue and node capacity for stable throughput.

During my internship at MIO Partners, I migrated ETL workflows from ActiveBatch/Talend to AWS Step Functions + S3 and built an internal portal exposing workflow logs, trigger history, and one-click execution controls to reduce developer triage and run-ops time.

Beyond enterprise work, I built Media Magic as a developer-first image delivery platform on AWS (S3, SQS, Lambda, DynamoDB) with SDK/CLI-based integration. It delivers sub-200ms responses, uses OpenCV + YOLO smart-focus cropping, applies idempotency and request coalescing to prevent duplicate transforms, and serves WebP variants through CloudFront with reliability-aware retries and delivery analytics.

I also developed Forest Fire Tracker, a Kafka-driven streaming platform that processes Unity-generated satellite image grids, applies OpenCV-based fire-region extraction and spread estimation, and provides live operational visibility through Redis-backed state management, Dask aggregation, and a Dash dashboard.

I am passionate about building robust, cloud-native data pipelines that handle both high-throughput streaming and batch workloads, and I am eager to bring this expertise to [Company] to solve complex data challenges.

Thank you for your time and consideration. I look forward to the opportunity to discuss how my skills and experience can contribute to your team.

Sincerely,
Umang Shah